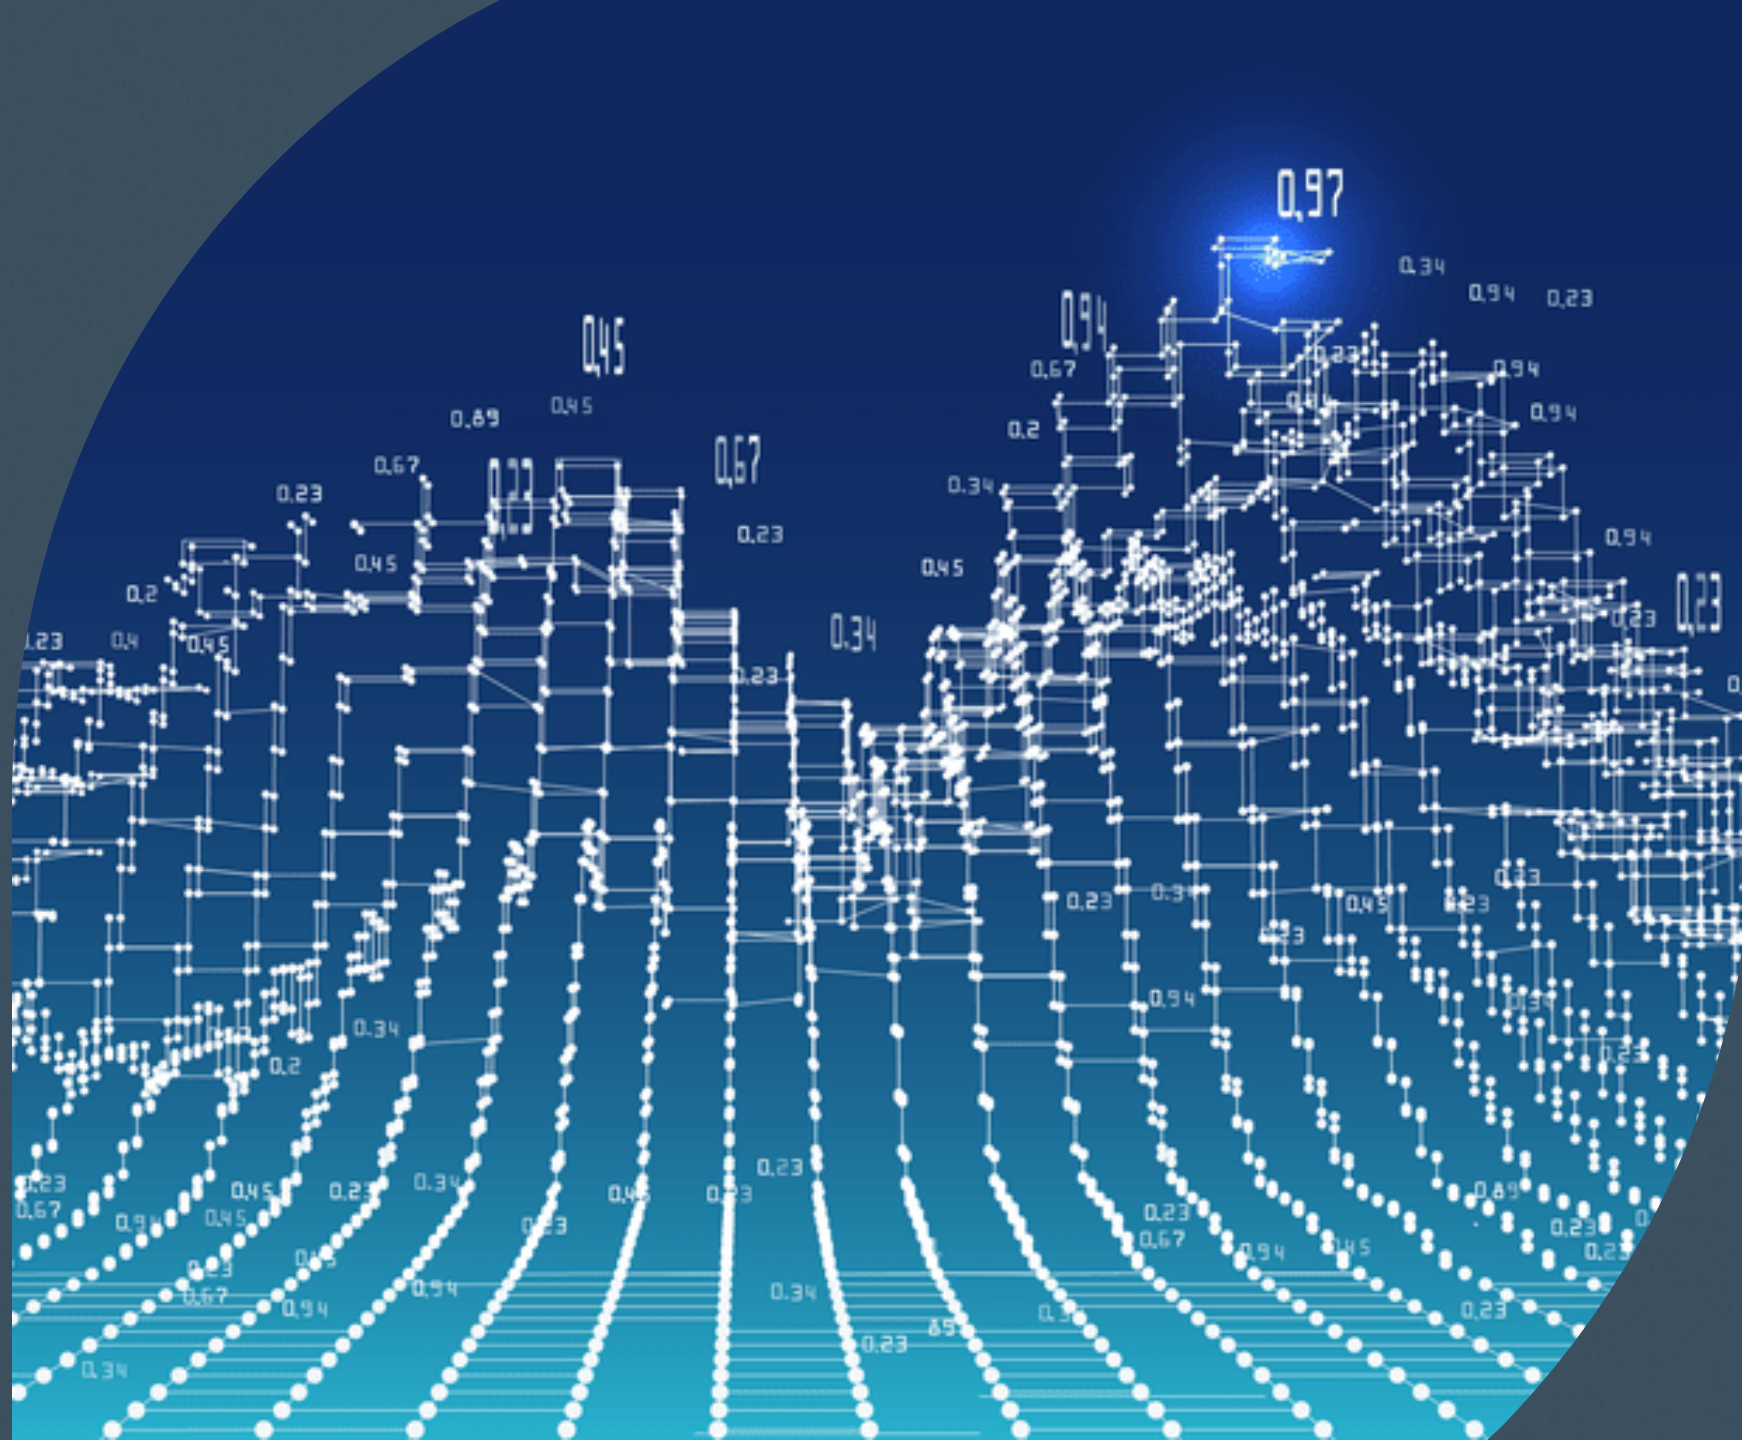


Quality Control

October 2024



Our Case Study

- A car manufacturer aims to enhance and automate the inspection of critical vehicle labels during the production process.
- Different label type:
 - vehicle identification number (VIN)
 - tire pressure labels
- Inspection Process :
 - Photographed using stationary cameras and mobile devices



Challenge

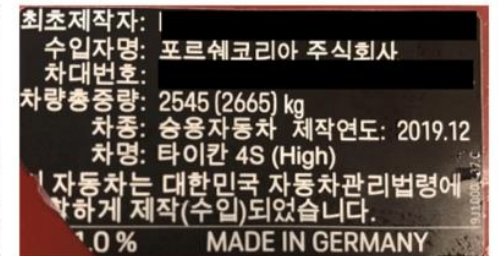

Varying lighting environment

Camera angles and distances

Use mobile camera for taking picture

Large variety of training data

- Labels differ by country with no consistent structure



Supervised Learning

labeled instances
are rarely available

Impossible to
collect sufficient
number of labeled

Costly to
synthetically
samples

Often dont known
which types of
defects may happen

Unsupervised Learning

Train only on
normal instances.

Learn normal
quality then
deviations can be
detected

Not relying on pre-
defined classes and
labeled samples

Dataset

2703 undamaged images

970 damaged images

Average resolution of the original images
was 1347×723

Images were cropped out, standardized,
and reduced to a unified size.

All images resized to 256×256 pixels

Use Mask R-CNN to segment labels

Approaches

01

Skip-
GANomaly(unsupervised)

02

PaDiM(unsupervised)

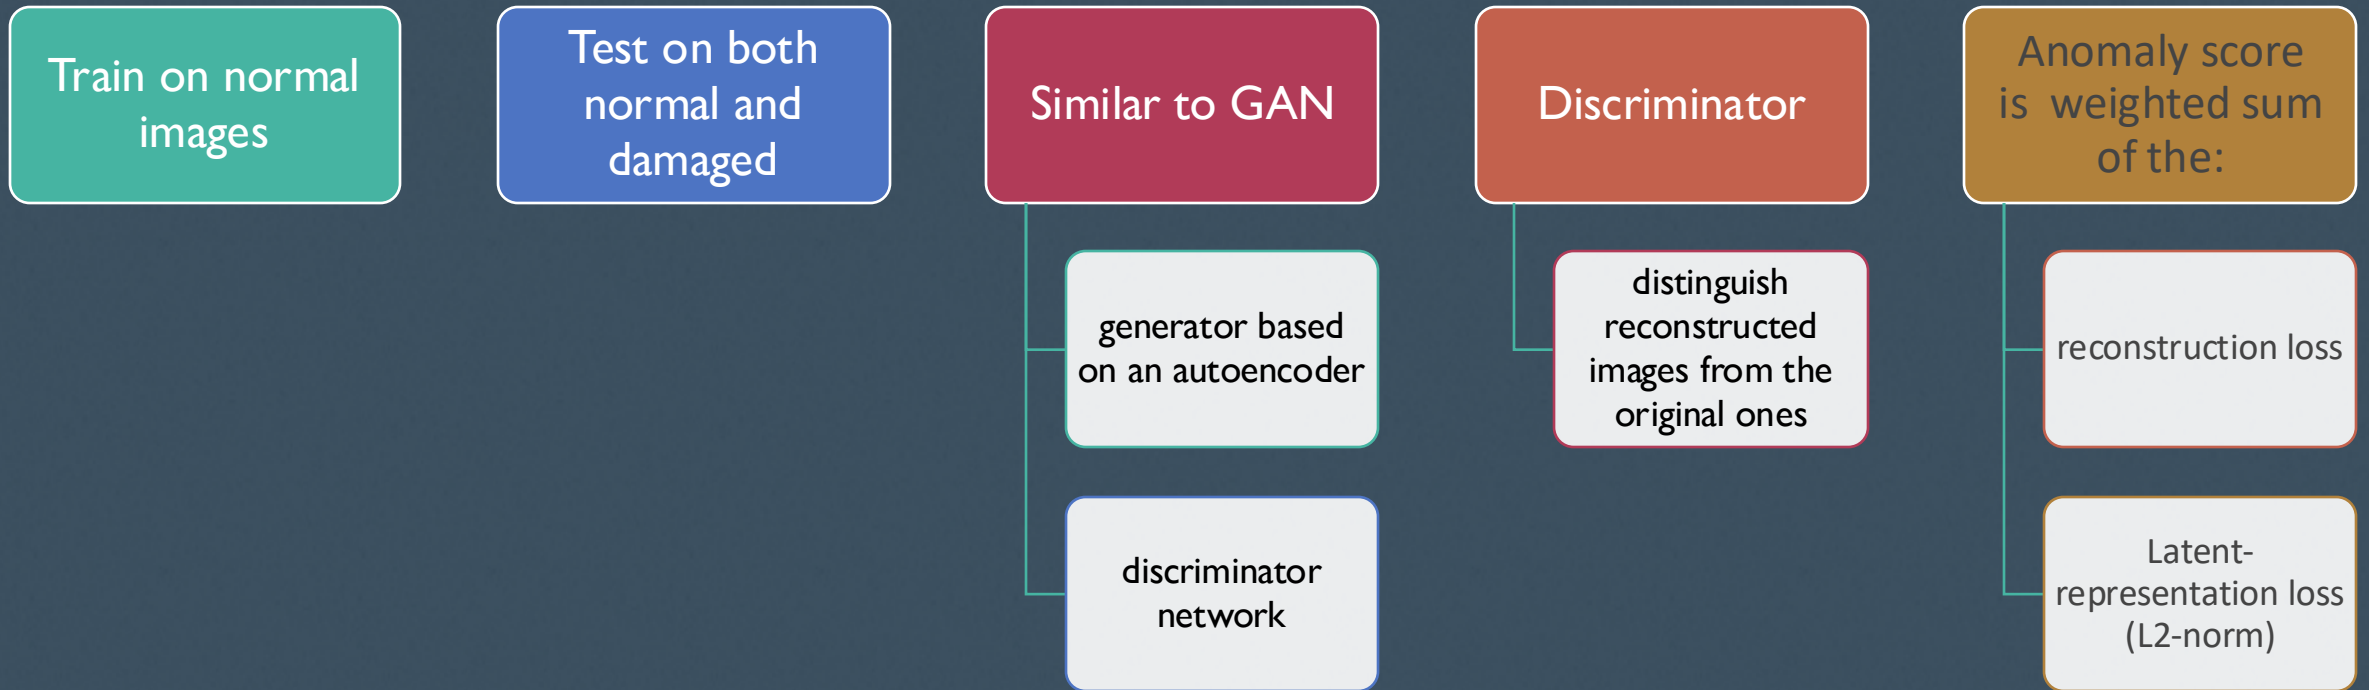
03

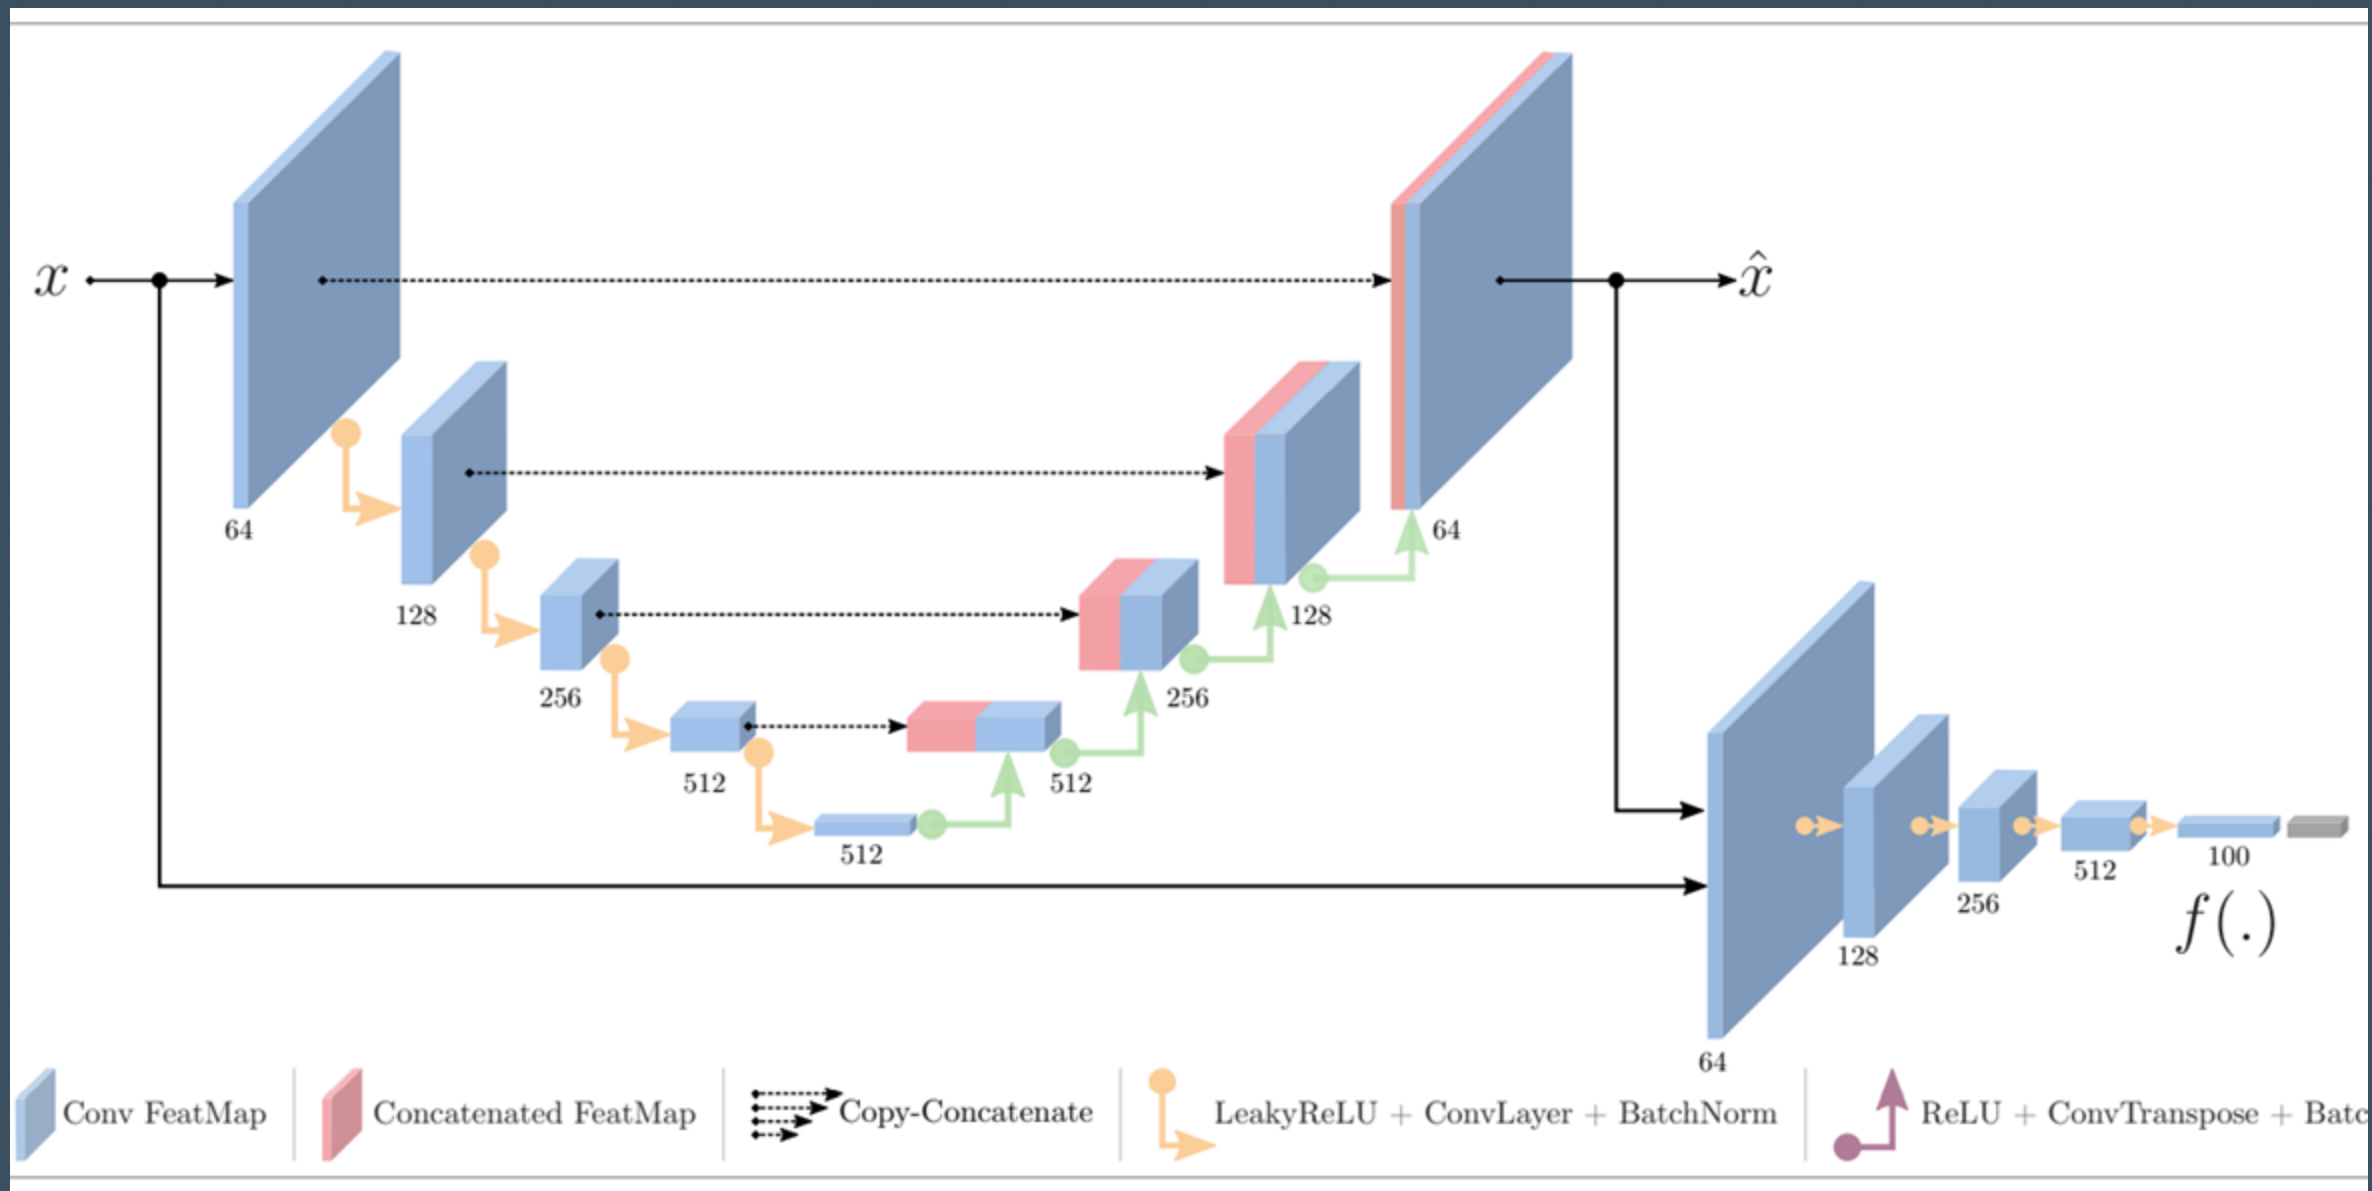
PatchCore(unsupervised)

04

Auto-Classifier(supervised)

Skip-GANomaly

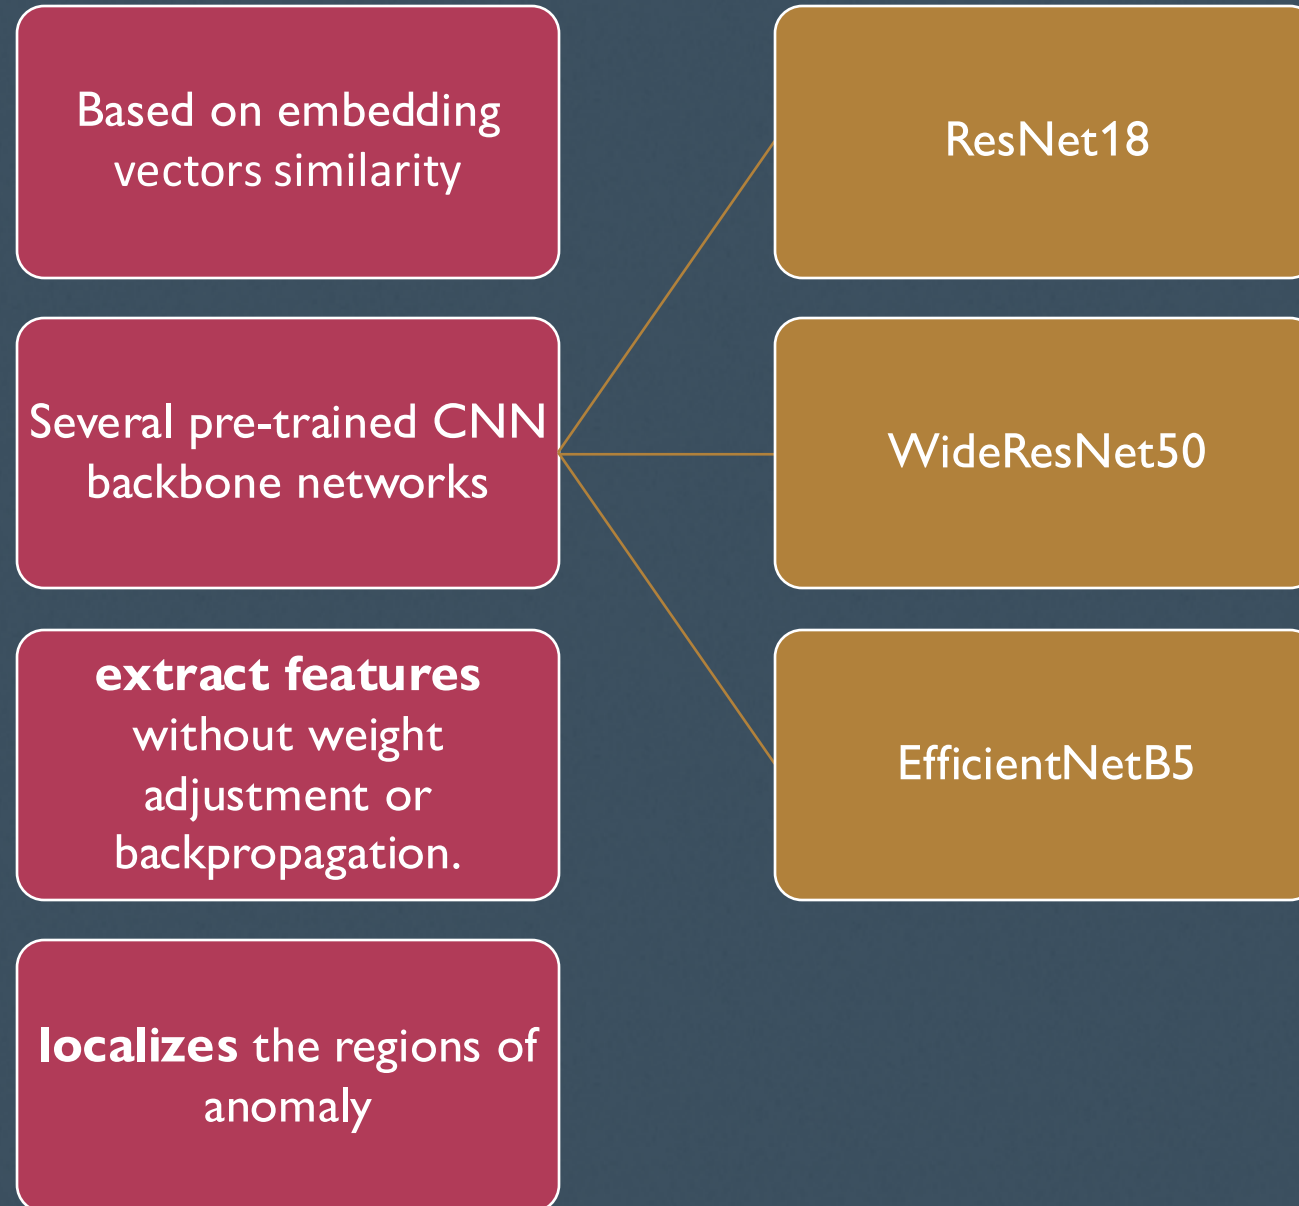





LOSS FUNCTIONS

- Use 3 different loss functions at the same time:
 - Adversarial Loss:
 - ensures that the network G reconstructs image as realistically as possible.
 - Contextual Loss:
 - explicitly learn contextual information
 - Latent Loss:
 - reconstruct latent representations for the input x and generated image as similar as possible

PaDiM





Other information

To reduce size of embedding vectors, randomly selecting few dimensions

More efficient than PCA

Anomaly score is distance between the test patch embedding x_{ij} and learned distribution

$$M(x_{ij}) = \sqrt{(x_{ij} - \mu_{ij})^T \Sigma_{ij}^{-1} (x_{ij} - \mu_{ij})}$$

PatchCore

Only requires normal images for training

Like PaDiM, extract embedding vectors

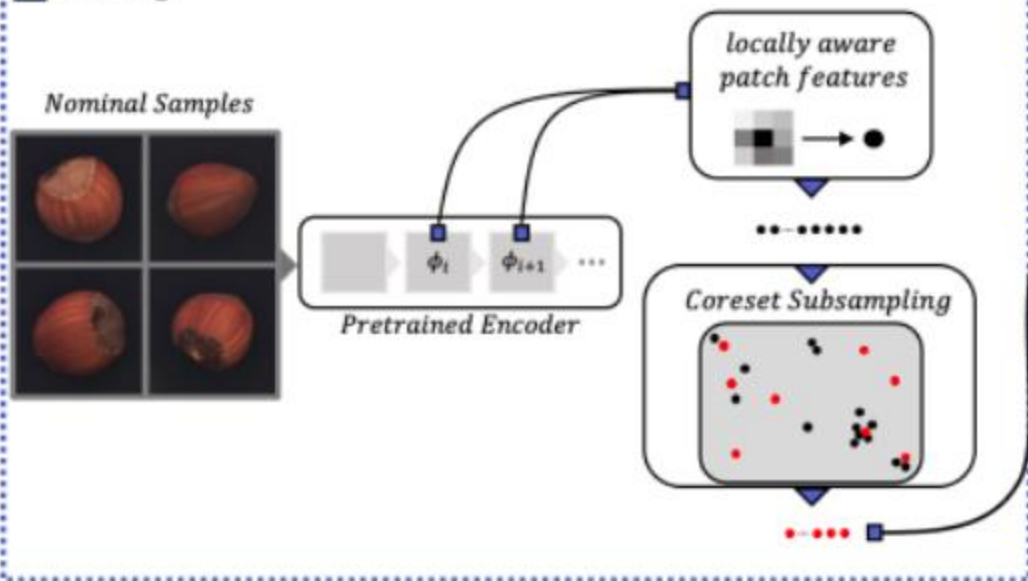
Use ImageNet

localizes the regions of anomaly


Process:

- First, embedding vectors stored in a memory bank
- To reduce memory bank size, subsampled by applying a k -center-greedy algorithm
 - sampling only 1% of the patch representations to be in the memory bank is sufficient to get good performance

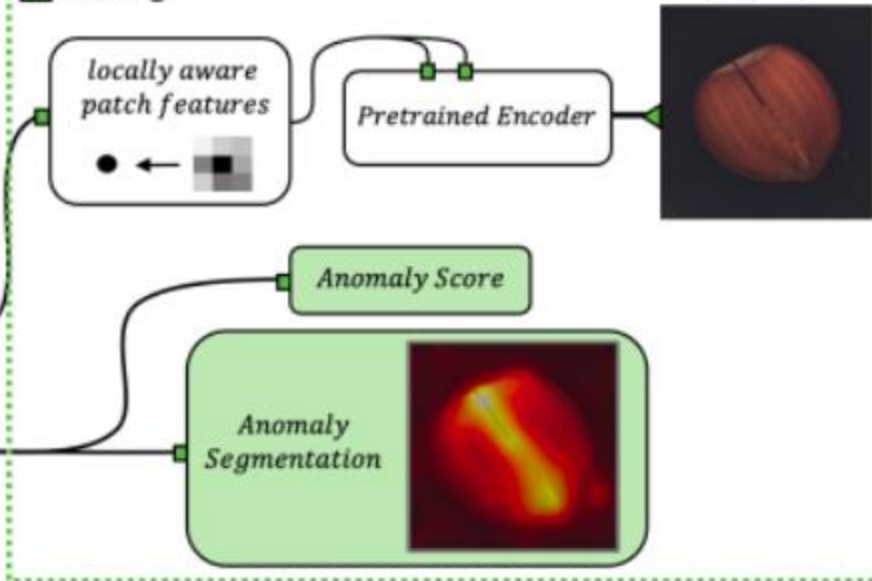
Training



PatchCore



Testing




Testing

Extract embedding
vectors of the input
image in the first step

Calculate distances to
the vectors of the
embedding
coreset using the k -
nearest neighbor

anomaly score is
calculated based on
nearest neighbor
weighted by the other k
distances

Example



Auto-Classfier

- A supervised approach
- Works on 2D images
- Balance the importance of each network through the process of normalization
- Networks with higher AUC scores have higher confidence

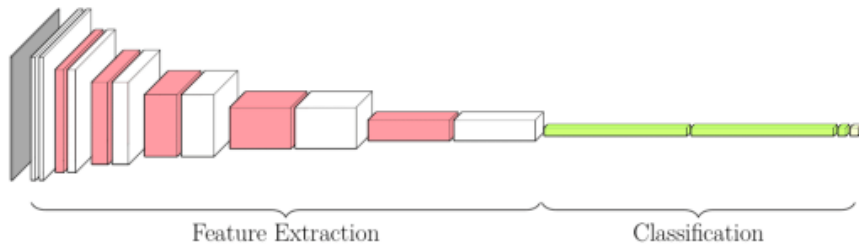
Process:

- Consider several CNN architectures
 - VGG11
 - VGG16
 - VGG19
 - ResNet18
- Fusing all the individual predictions into a final, weighted, prediction by making a weighted sum of each class


$$w_i = \frac{V_i}{\sum_{j=1}^n V_j}, i \in 1, \dots, n$$

Auto-Classfier

Second part:



(a) Individual CNN



(b) Auto-Classfier

Results

Table 3

Evaluation results of the anomaly detection performance based on the test set.

Model	Backbone	<i>AUROC</i>	<i>F1 Score</i>	<i>Recall</i> _{$p_r=0.996$}	<i>Recall</i> _{$p_r=0.95$}	<i>Recall</i> _{$p_r=0.9$}
Auto-Classifier	ResNet50	1.000 ± 0.000	0.995 ± 0.005	0.995 ± 0.009	1.000 ± 0.000	1.000 ± 0.000
	Fusion	1.000 ± 0.001	0.995 ± 0.006	0.792 ± 0.443	1.000 ± 0.000	1.000 ± 0.000
	AutoML	1.000 ± 0.000	0.992 ± 0.003	0.978 ± 0.022	1.000 ± 0.000	1.000 ± 0.000
Skip-GANomaly	–	0.930 ± 0.006	0.757 ± 0.009	0.130 ± 0.081	0.214 ± 0.166	0.275 ± 0.160
PaDiM	EfficientNetB5	0.992 ± 0.002	0.911 ± 0.009	0.438 ± 0.240	0.884 ± 0.031	0.940 ± 0.020
	ResNet18	0.978 ± 0.003	0.874 ± 0.010	0.230 ± 0.092	0.486 ± 0.206	0.790 ± 0.048
	ResNet50	0.982 ± 0.005	0.897 ± 0.017	0.205 ± 0.143	0.697 ± 0.112	0.841 ± 0.080
	WideResNet50	0.984 ± 0.005	0.893 ± 0.033	0.318 ± 0.113	0.712 ± 0.136	0.867 ± 0.087
PatchCore	WideResNet50	0.996 ± 0.002	0.948 ± 0.015	0.451 ± 0.245	0.941 ± 0.033	0.986 ± 0.017

Training time

Table 4

Evaluation results of training and inference times for a single run.

Model	Backbone	Train time [h]	Inference time [ms]
Auto-Classifier	ResNet50	1.401	17.48
	Fusion	15.528	199.96
	AutoML	2.433	10.65
Skip-GANomaly	–	4.935	150.39
PaDiM	ResNet18	0.211	57.89
	ResNet50	0.236	111.98
	WideResNet50	0.241	119.96
	EfficientNetB5	0.248	103.62
PatchCore	WideResNet50	44.491	1818.52